

Road Accidents Analysis & Severity Prediction

Importing necessary Libraries

```
In [4]: %%time
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import warnings
import requests
import seaborn as sns
import plotly.express as px
import plotly.graph_objects as go

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

CPU times: total: 3.38 s
Wall time: 8.49 s
```


Loading Dataset

```
In [5]: data = pd.read_excel('RTAdata.xlsx')
data.head()
```

Out[5]:

	Time	Day_of_week	Age_band_of_driver	Sex_of_driver	Educational_level	Vehicle_driver_relation	Driving_experience	Type_of_vehicle	Owner_of_vehicle
0	17:02:00	Monday	18-30	Male	Above high school	Employee	1-2yr	Automobile	Owner
1	17:02:00	Monday	31-50	Male	Junior high school	Employee	Above 10yr	Public (> 45 seats)	Owner
2	17:02:00	Monday	18-30	Male	Junior high school	Employee	1-2yr	Lorry (41-100Q)	Owner
3	01:06:00	Sunday	18-30	Male	Junior high school	Employee	5-10yr	Public (> 45 seats)	Governmental
4	01:06:00	Sunday	18-30	Male	Junior high school	Employee	2-5yr	NaN	Owner

5 rows × 32 columns



```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12316 entries, 0 to 12315
Data columns (total 32 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Time                12316 non-null object
1   Day_of_week         12316 non-null object
2   Age_band_of_driver  12316 non-null object
3   Sex_of_driver       12316 non-null object
4   Educational_level    11575 non-null object
```

Cleaning Dataset

```
In [7]: # Remove rows with "Unknown" values
df_cleaned = data[~data.apply(lambda row: row.astype(str).str.contains('Unknown').any(), axis=1)]
cleaned_file_path = 'cleaned_dataset.xlsx'
df_cleaned.to_excel(cleaned_file_path, index=False)
print(f"Cleaned data saved to {cleaned_file_path}")
```

Cleaned data saved to cleaned_dataset.xlsx

```
In [8]: data1 = pd.read_excel('cleaned_dataset.xlsx')
data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7417 entries, 0 to 7416
Data columns (total 32 columns):
```

```
In [10]: data1.describe().T
```

```
Out[10]:
```

	count	mean	std	min	25%	50%	75%	max
Number_of_vehicles_involved	7417.0	2.036807	0.668984	1.0	2.0	2.0	2.0	7.0
Number_of_casualties	7417.0	1.586086	1.042470	1.0	1.0	1.0	2.0	8.0
Casualty_severity	4675.0	2.898810	0.313694	1.0	3.0	3.0	3.0	3.0

```
In [11]: #skewness
selected_columns1=["Number_of_vehicles_involved","Number_of_casualties"]
skewness = data1[selected_columns1].skew()
print(skewness)
```

```
Number_of_vehicles_involved    1.260581
Number_of_casualties           2.227987
dtype: float64
```

```
In [12]: for i in data1.columns:
          if data1[i].dtypes== object:
              print(i)
              print(data1[i].unique())
              print(data1[i].nunique())
              print()S
```

```
Day_of_week
['Monday' 'Sunday' 'Wednesday' 'Friday' 'Saturday' 'Thursday' 'Tuesday']
7

Age_band_of_driver
['18-30' '31-50' 'Under 18' 'Over 51']
4

Sex_of_driver
['Male' 'Female']
2

Educational_level
['Above high school' 'Junior high school' 'Elementary school' nan
 'High school' 'Illiterate' 'Writing & reading']
6

Vehicle_driver_relation
['Employee' 'Owner' nan 'Other']
3

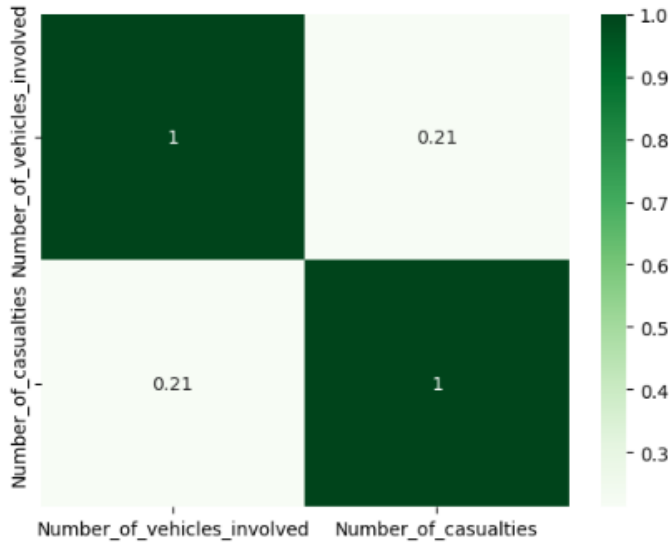
Driving_experience
['1-2yr' 'Above 10yr' '5-10yr' '2-5yr' 'No Licence' 'Below 1yr' nan
 'unknown']
7

Type_of_vehicle
['Automobile' 'Public (> 45 seats)' 'Lorry (41-100Q)' nan 'Lorry (11-40Q)'
 'Public (13-45 seats)' 'Long lorry' 'Public (12 seats)' 'Taxi'
 'Pick up upto 10Q' 'Stationwagen' 'Other' 'Ridden horse' 'Motorcycle'
 'Turbo' 'Bicycle' 'Special vehicle' 'Bajaj']
17
```

Correlation

```
In [15]: #for Numerical Variables
correlation = data[selected_columns1].corr()
sns.heatmap(correlation,annot = True, cmap = 'Greens')
```

Out[15]: <Axes: >



Cramer's V for Categorical Variables

```
In [73]: #for categorical variables
from scipy.stats import chi2_contingency
selected_columns2=["Age_band_of_casualty","Sex_of_casualty","Type_of_vehicle","Road_surface_type","Road_surface_conditions",
                  "Light_conditions","Type_of_collision","Accident_severity"]
data2 = data1[selected_columns2]

# Function to calculate Cramér's V
def cramers_v(confusion_matrix):
    chi2 = chi2_contingency(confusion_matrix)[0]
    n = confusion_matrix.sum().sum()
    r, k = confusion_matrix.shape
    return np.sqrt(chi2 / (n * (min(r, k) - 1)))

# Calculate Cramér's V for each variable with respect to the target variable
target_variable = 'Accident_severity'
results = {}
for variable in data2.columns:
    if variable != target_variable:
        contingency_table = pd.crosstab(data2[variable], data2[target_variable])
        cramers_v_value = cramers_v(contingency_table)
        results[variable] = cramers_v_value

# Display the results
for variable, value in results.items():
    print(f"Cramér's V for {variable} with respect to {target_variable}: {value}")
```


Cramér's V for Age_band_of_casualty with respect to Accident_severity: 0.04100781145890263
Cramér's V for Sex_of_casualty with respect to Accident_severity: 0.016888876631200308
Cramér's V for Type_of_vehicle with respect to Accident_severity: 0.060166950245805886
Cramér's V for Road_surface_type with respect to Accident_severity: 0.022444251595571993
Cramér's V for Road_surface_conditions with respect to Accident_severity: 0.018750916611646087
Cramér's V for Light_conditions with respect to Accident_severity: 0.0554158069505708
Cramér's V for Type_of_collision with respect to Accident_severity: 0.035515244289020614

Exploratory Data Analysis (EDA)

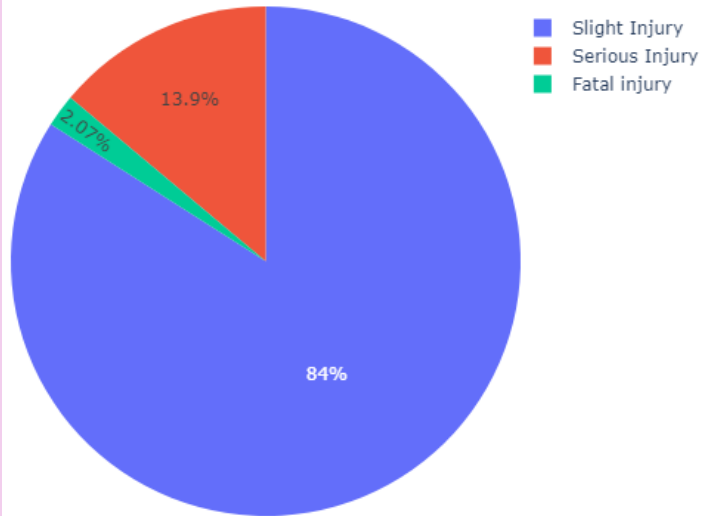
Pie Charts for Gender Distribution, Injury Type and Cause of Accidents

```
In [18]: #PIE Charts
fig1=px.pie(data1, data1['Sex_of_driver'], data1['Number_of_casualties'],title='Gender Distribution' ,template='seaborn')
fig2=px.pie(data1, data1['Accident_severity'], data1['Number_of_casualties'],title='Injury Type Distribution',template='plotly')
fig1.show()
fig2.show()
fig3=px.pie(data1,data1['Cause_of_accident'],data1['Number_of_casualties'],
            color='Cause_of_accident',template='ggplot2',title='Casualties by Cause of Accident')
fig3.show()
```


Gender Distribution



Injury Type Distribution

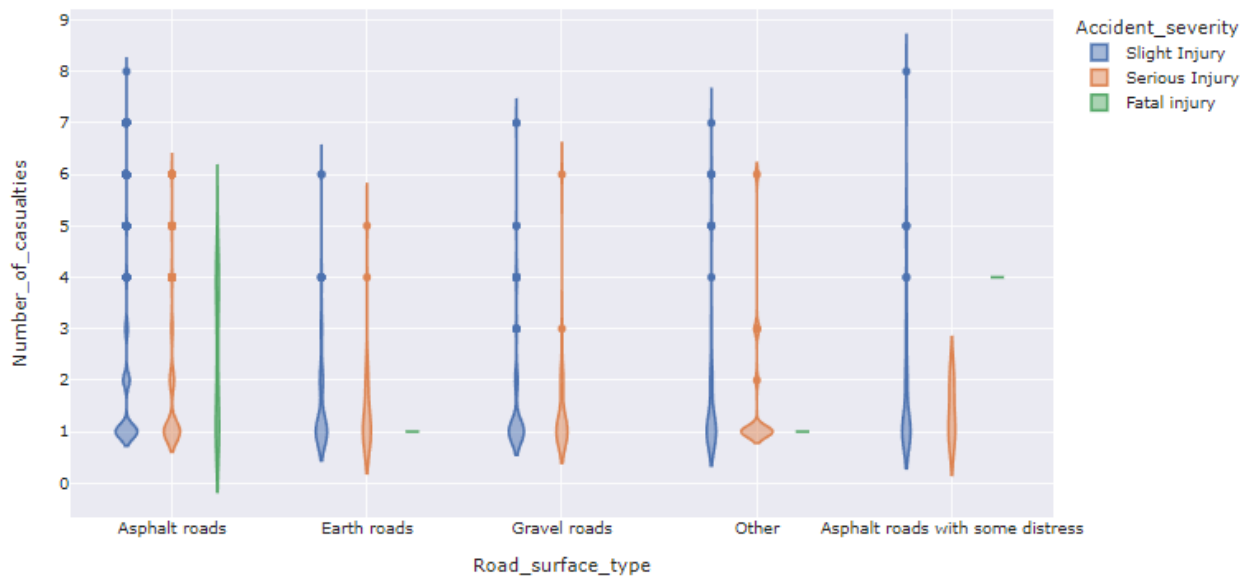


Casualties by Cause of Accident



Violin plot for Casualties by Road-surface type

```
In [19]: fig4=px.violin(data1,data1['Road_surface_type'], data1['Number_of_casualties'],color='Accident_severity',
template='seaborn')
fig4.show()
```




Bar Charts for Casualties by Day of the week, Education-level

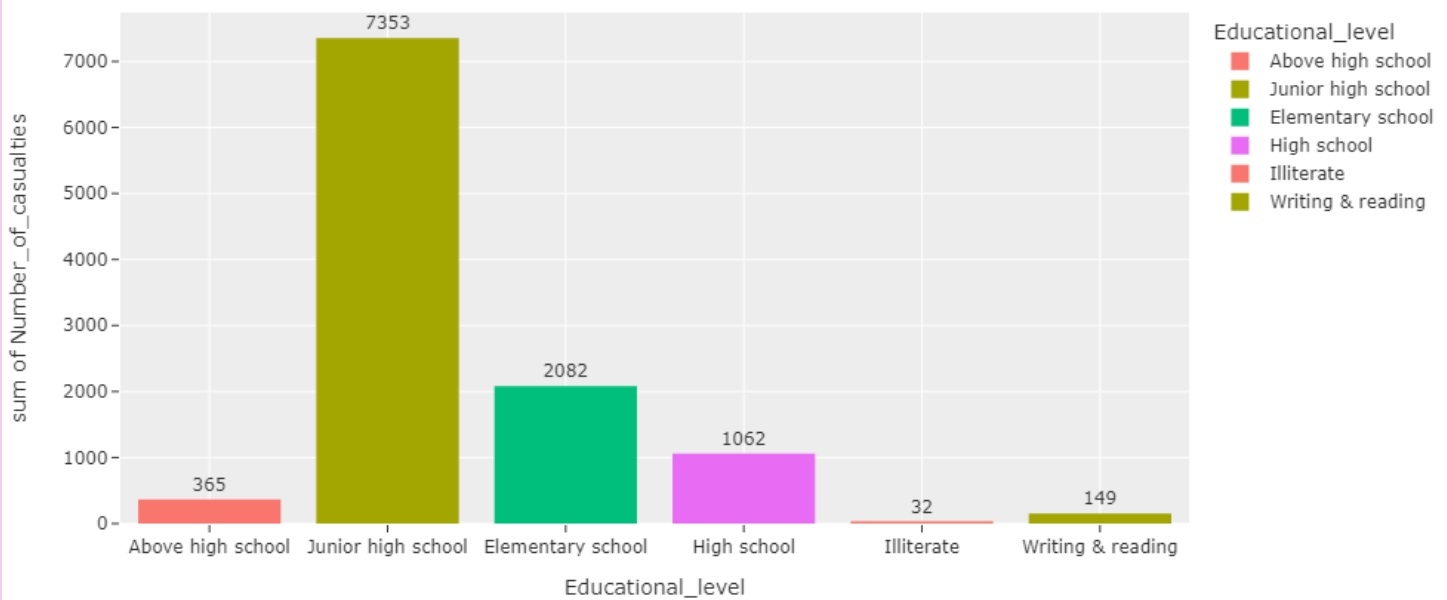
```
In [20]: fig5=px.histogram(data1,data1['Day_of_week'],data1['Number_of_casualties'],
color='Day_of_week',template='seaborn',title='Casualties by Day of the week',
labels={'Number_of_casualties': 'Number of Casualties'})
fig5.update_traces(texttemplate='%{y}', textposition='outside')

fig6=px.histogram(data1,data1['Educational_level'],data1['Number_of_casualties'],
color='Educational_level',template='ggplot2',title='Casualties by Educational Level')
fig6.update_traces(texttemplate='%{y}', textposition='outside')
fig5.show()
fig6.show()
```

Casualties by Day of the week




Casualties by Educational Level



Bar diagram for Casualties by Time of the day

```
In [21]: fig7=px.histogram(data1,data1['Time'],data1['Number_of_casualties']
,template='ggplot2',title='Casualties by Time of the day')
fig7.update_traces(texttemplate='%{y}', textposition='outside')
fig7.show()
```

Casualties by Time of the day



Bar diagram above doesn't help us , so we will create 3-hour interval groups.

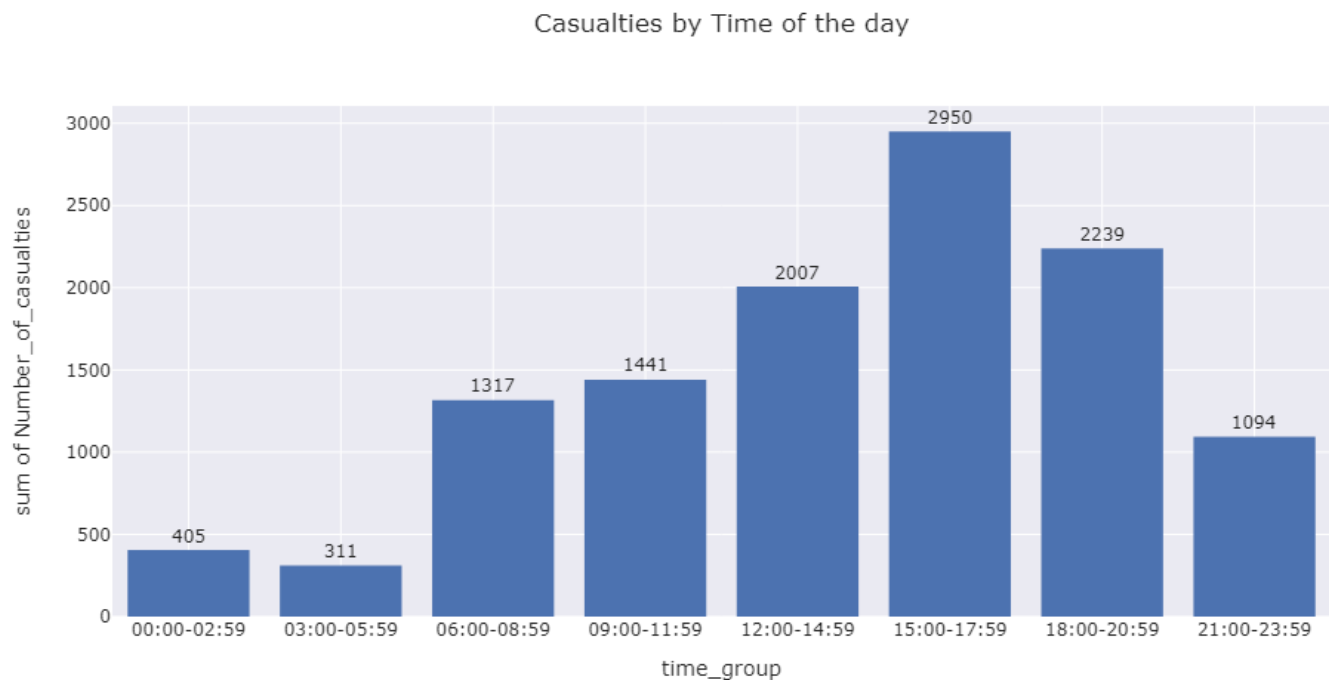
Bar diagram above doesn't help us, so we will create 3-hour interval groups

```
In [11]: print(data1['Time'].dtypes)
# Convert time to 3-hour interval groups
bins = pd.to_timedelta(['00:00:00', '03:00:00', '06:00:00', '09:00:00',
                        '12:00:00', '15:00:00', '18:00:00', '21:00:00', '23:59:59'])
labels = ['00:00-02:59', '03:00-05:59', '06:00-08:59', '09:00-11:59',
          '12:00-14:59', '15:00-17:59', '18:00-20:59', '21:00-23:59']
data1['time_group'] = pd.cut(pd.to_datetime(data1['Time']),
                             format='%H:%M:%S').dt.time.apply(lambda x: pd.to_timedelta(str(x)),
                             bins=bins, labels=labels)
print(data1[['Time', 'time_group']])
```

```
object
      Time  time_group
0  17:02:00  15:00-17:59
1  17:02:00  15:00-17:59
2  17:02:00  15:00-17:59
3   01:06:00  00:00-02:59
4   01:06:00  00:00-02:59
...      ...      ...
```

Bar diagram for Casualties by Time of the day

```
In [ ]: # Define the order of time groups
time_group_order = ['00:00-02:59', '03:00-05:59', '06:00-08:59', '09:00-11:59',
                    '12:00-14:59', '15:00-17:59', '18:00-20:59', '21:00-23:59']
fig8=px.histogram(data1,data1['time_group'],data1['Number_of_casualties'],
                  template='seaborn',title='Casualties by Time of the day',category_orders={'time_group': time_group_order})
fig8.update_traces(texttemplate='%{y}', textposition='outside')
fig8.show()
```



Data Pre-processing

```
In [13]: #making a duplicate dataset
data3=data1.copy(deep=True)
data3.head()
```

```
Out[13]:
```

ualty	Age_band_of_casualty	Casualty_severity	Work_of_casualty	Fitness_of_casualty	Pedestrian_movement	Cause_of_accident	Accident_severity	time_group
NaN	NaN	NaN	NaN	NaN	Not a Pedestrian	Moving Backward	Slight Injury	15:00-17:59
NaN	NaN	NaN	NaN	NaN	Not a Pedestrian	Overtaking	Slight Injury	15:00-17:59
Male	31-50	3.0	Driver	NaN	Not a Pedestrian	Changing lane to the left	Serious Injury	15:00-17:59
Female	18-30	3.0	Driver	Normal	Not a Pedestrian	Changing lane to the right	Slight Injury	00:00-02:59
NaN	NaN	NaN	NaN	NaN	Not a Pedestrian	Overtaking	Slight Injury	00:00-02:59

Selecting required columns

```
In [33]: columns_for_model=["time_group","Accident_severity","Day_of_week","Type_of_vehicle","Road_surface_type","Type_of_collision","Casualty_class","Sex_of_casualty"]
data_for_model=data3[columns_for_model]
data_for_model.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7417 entries, 0 to 7416
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   time_group            7417 non-null   category
1   Accident_severity     7417 non-null   object
2   Day_of_week           7417 non-null   object
3   Type_of_vehicle       6831 non-null   object
4   Road_surface_type     7314 non-null   object
5   Type_of_collision     7326 non-null   object
6   Casualty_class        4675 non-null   object
7   Sex_of_casualty       4675 non-null   object
dtypes: category(1), object(7)
```

```
In [34]: data_for_model.head()
```

```
Out[34]:
```

	time_group	Accident_severity	Day_of_week	Type_of_vehicle	Road_surface_type	Type_of_collision	Casualty_class	Sex_of_casualty
0	15:00-17:59	Slight Injury	Monday	Automobile	Asphalt roads	Collision with roadside-parked vehicles	NaN	NaN
1	15:00-17:59	Slight Injury	Monday	Public (> 45 seats)	Asphalt roads	Vehicle with vehicle collision	NaN	NaN
2	15:00-17:59	Serious Injury	Monday	Lorry (41-100Q)	Asphalt roads	Collision with roadside objects	Driver or rider	Male
3	00:00-02:59	Slight Injury	Sunday	Public (> 45 seats)	Earth roads	Vehicle with vehicle collision	Pedestrian	Female
4	00:00-02:59	Slight Injury	Sunday	NaN	Asphalt roads	Vehicle with vehicle collision	NaN	NaN

Since data type is object, we need to convert them into nominal type before applying any machine learning technique.

```
In [35]: #converting object type values to nominal values.
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for col in data_for_model.columns:
    if data_for_model[col].dtype == object or data_for_model[col].dtype.name == 'category':
        data_for_model[col] = le.fit_transform(data_for_model[col])
```

```
In [37]: data_RF=data_for_model.copy(deep=True)
data_RF.head()
data_RF.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7417 entries, 0 to 7416
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   time_group            7417 non-null   int32
1   Accident_severity     7417 non-null   int32
2   Day_of_week           7417 non-null   int32
3   Type_of_vehicle       7417 non-null   int32
4   Road_surface_type     7417 non-null   int32
5   Type_of_collision     7417 non-null   int32
6   Casualty_class        7417 non-null   int32
7   Sex_of_casualty       7417 non-null   int32
dtypes: int32(8)
```

Random-Forest Classifier

```
In [57]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score, recall_score, f1_score, precision_score

x1=data_RF.drop('Accident_severity',axis=1)
y1=data_RF['Accident_severity']
xtrain, xtest, ytrain, ytest = train_test_split(x1,y1,test_size=0.30)
rf = RandomForestClassifier()
rf.fit(xtrain,ytrain)

ypred=rf.predict(xtest)

print('confusion matrix : ',confusion_matrix(ytest,ypred))
print('classification report:',classification_report(ytest,ypred))
print('accuracy : ',round(accuracy_score(ytest,ypred),2))
print('precision : ',round(precision_score(ytest,ypred,average='weighted'),2))
print('recall : ',round(recall_score(ytest,ypred,average='weighted'),2))
print('f1 : ',round(f1_score(ytest,ypred,average='weighted'),2))
print()
```

```
confusion matrix : [[ 2   1  23]
 [ 1 14 294]
 [ 1 83 1807]]
classification report:
              precision    recall  f1-score   support

    0       0.50      0.08   0.13         26
    1       0.14      0.05   0.07        309
    2       0.85      0.96   0.90       1891

 accuracy          0.82      2226
macro avg          0.50      0.36   0.37      2226
weighted avg       0.75      0.82   0.78      2226

accuracy : 0.82
precision : 0.75
recall : 0.82
f1 : 0.78
```

```
In [58]: #variance importance and Gini coefficient
feature_importances = rf.feature_importances_
features = x1.columns
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': feature_importances
})

importance_df = importance_df.sort_values(by='Importance', ascending=False)

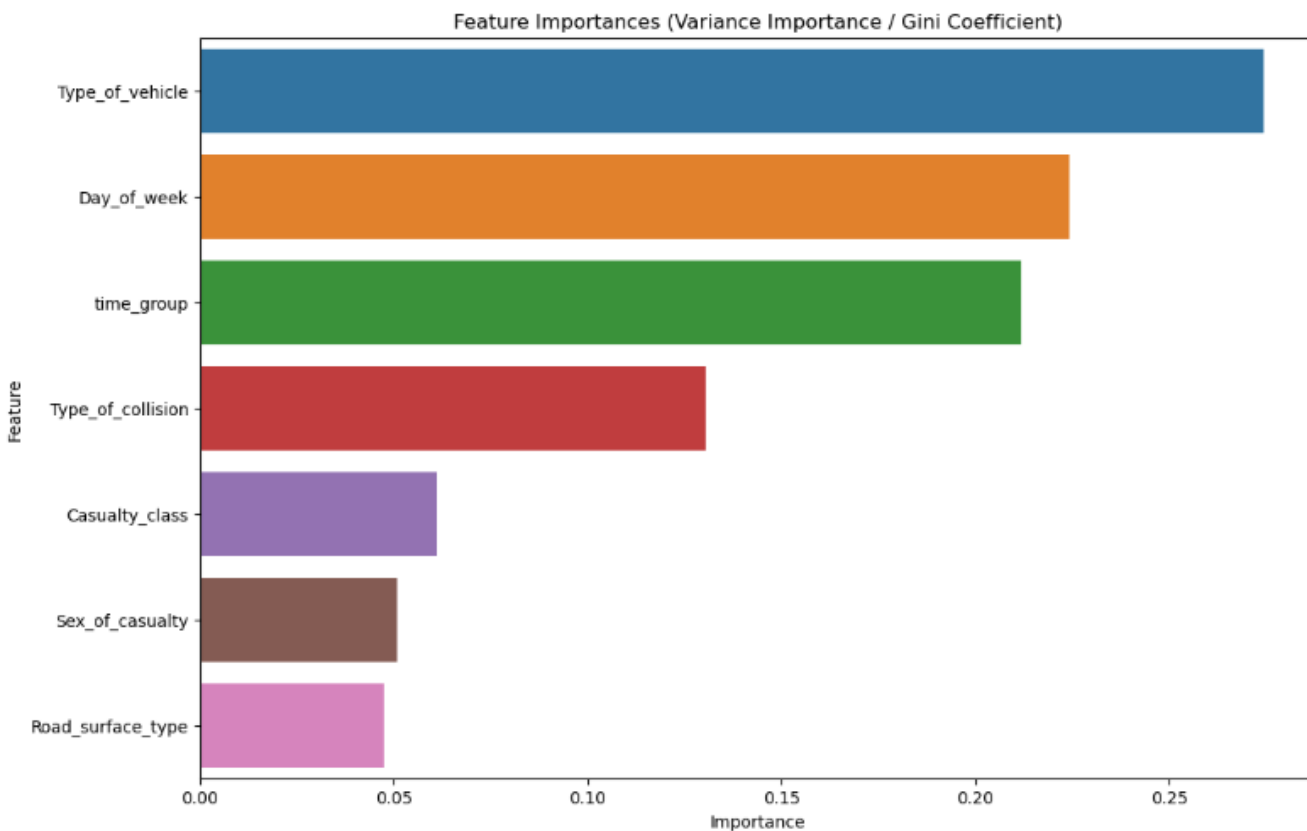
print('Feature Importances (Variance Importance / Gini Coefficient):')
print(importance_df)

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=importance_df)
plt.title('Feature Importances (Variance Importance / Gini Coefficient)')
plt.show()
```

Feature Importances (Variance Importance / Gini Coefficient):

	Feature	Importance
2	Type_of_vehicle	0.274236
1	Day_of_week	0.224294
0	time_group	0.211873
4	Type_of_collision	0.130441
5	Casualty_class	0.061067
6	Sex_of_casualty	0.050696
3	Road_surface_type	0.047393



K-Nearest Neighbour

```
In [41]: data_KNN=data_for_model.copy(deep=True)
data_KNN.head()
data_KNN.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7417 entries, 0 to 7416
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   time_group             7417 non-null   int32  
1   Accident_severity      7417 non-null   int32  
2   Day_of_week            7417 non-null   int32  
3   Type_of_vehicle        7417 non-null   int32  
4   Road_surface_type      7417 non-null   int32  
5   Type_of_collision      7417 non-null   int32  
6   Casualty_class         7417 non-null   int32  
7   Sex_of_casualty        7417 non-null   int32  
dtypes: int32(8)
memory usage: 231.9 KB
```

```
In [61]: from sklearn.neighbors import KNeighborsClassifier
X2 = data_KNN.drop('Accident_severity', axis=1)
y2 = data_KNN['Accident_severity']

X_train, X_test, y_train, y_test = train_test_split(X2, y2, test_size=0.2)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
```

```
In [53]: y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')

cm = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(cm)

print('Classification Report:')
print(classification_report(y_test, y_pred))

Accuracy: 0.8140161725067385
Confusion Matrix:
[[ 0  1 13]
 [ 0  5 222]
 [ 4 36 1203]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	14
1	0.12	0.02	0.04	227
2	0.84	0.97	0.90	1243
accuracy			0.81	1484
macro avg	0.32	0.33	0.31	1484
weighted avg	0.72	0.81	0.76	1484

Permutation Importance

```
In [55]: from sklearn.inspection import permutation_importance
perm_importance = permutation_importance(knn, X_test, y_test, n_repeats=30, random_state=42)

features = data_KNN.drop('Accident_severity', axis=1).columns
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': perm_importance.importances_mean,
    'Standard Deviation': perm_importance.importances_std})

importance_df = importance_df.sort_values(by='Importance', ascending=False)

print('Permutation Feature Importances:')
print(importance_df)
```

```
Permutation Feature Importances:
   Feature  Importance  Standard Deviation
4  Type_of_collision    0.002583         0.003127
5  Casualty_class      0.000022         0.003758
0    time_group      -0.000517         0.003012
6  Sex_of_casualty     -0.000674         0.002723
1    Day_of_week     -0.002156         0.004174
2  Type_of_vehicle     -0.002673         0.004492
3  Road_surface_type   -0.003077         0.001711
```

XGBOOST(eXtreme Gradient Boosting)

```
In [62]: data_XGB=data_for_model.copy(deep=True)
data_XGB.head()
data_XGB.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7417 entries, 0 to 7416
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   time_group            7417 non-null   int32
1   Accident_severity     7417 non-null   int32
2   Day_of_week           7417 non-null   int32
3   Type_of_vehicle       7417 non-null   int32
4   Road_surface_type     7417 non-null   int32
5   Type_of_collision     7417 non-null   int32
6   Casualty_class        7417 non-null   int32
7   Sex_of_casualty       7417 non-null   int32
dtypes: int32(8)
memory usage: 231.9 KB
```

```
In [84]: import xgboost as xgb
X3 = data_XGB.drop('Accident_severity', axis=1)
y3 = data_XGB['Accident_severity']

X_train, X_test, y_train, y_test = train_test_split(X3, y3, test_size=0.2)

dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)

params = {'objective': 'multi:softmax', 'num_class': 3, 'eval_metric': 'mlogloss', 'eta': 0.3, 'max_depth':
        6, 'min_child_weight': 1,
        'subsample': 1, 'colsample_bytree': 1}
```

```
In [85]: num_rounds = 100

evals = [(dtrain, 'train'), (dtest, 'eval')]
model = xgb.train(params, dtrain, num_boost_round=num_rounds, evals=evals)

y_pred = model.predict(dtest)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')

print('Classification Report:')
print(classification_report(y_test, y_pred))
```

```
[0] train-mlogloss:0.86080 eval-mlogloss:0.86569
[1] train-mlogloss:0.72372 eval-mlogloss:0.73336
[2] train-mlogloss:0.63806 eval-mlogloss:0.65146
[3] train-mlogloss:0.58258 eval-mlogloss:0.59957
[4] train-mlogloss:0.54459 eval-mlogloss:0.56653
[5] train-mlogloss:0.51835 eval-mlogloss:0.54410
[6] train-mlogloss:0.49906 eval-mlogloss:0.52861
[7] train-mlogloss:0.48518 eval-mlogloss:0.51847
[8] train-mlogloss:0.47384 eval-mlogloss:0.51271
[9] train-mlogloss:0.46532 eval-mlogloss:0.50776
[10] train-mlogloss:0.45840 eval-mlogloss:0.50476
[11] train-mlogloss:0.45302 eval-mlogloss:0.50200
[12] train-mlogloss:0.44710 eval-mlogloss:0.50073
[13] train-mlogloss:0.44243 eval-mlogloss:0.49962
```

```

[73] train-mlogloss:0.32907 eval-mlogloss:0.53203
[74] train-mlogloss:0.32925 eval-mlogloss:0.53290
[75] train-mlogloss:0.32815 eval-mlogloss:0.53410
[76] train-mlogloss:0.32704 eval-mlogloss:0.53475
[77] train-mlogloss:0.32613 eval-mlogloss:0.53505
[78] train-mlogloss:0.32550 eval-mlogloss:0.53565
[79] train-mlogloss:0.32463 eval-mlogloss:0.53643
[80] train-mlogloss:0.32377 eval-mlogloss:0.53721
[81] train-mlogloss:0.32279 eval-mlogloss:0.53733
[82] train-mlogloss:0.32206 eval-mlogloss:0.53752
[83] train-mlogloss:0.32132 eval-mlogloss:0.53803
[84] train-mlogloss:0.32026 eval-mlogloss:0.53924
[85] train-mlogloss:0.31947 eval-mlogloss:0.54055
[86] train-mlogloss:0.31894 eval-mlogloss:0.54092
[87] train-mlogloss:0.31798 eval-mlogloss:0.54206
[88] train-mlogloss:0.31711 eval-mlogloss:0.54241
[89] train-mlogloss:0.31605 eval-mlogloss:0.54246
[90] train-mlogloss:0.31554 eval-mlogloss:0.54292
[91] train-mlogloss:0.31504 eval-mlogloss:0.54383
[92] train-mlogloss:0.31425 eval-mlogloss:0.54453
[93] train-mlogloss:0.31357 eval-mlogloss:0.54524
[94] train-mlogloss:0.31305 eval-mlogloss:0.54540
[95] train-mlogloss:0.31254 eval-mlogloss:0.54617
[96] train-mlogloss:0.31211 eval-mlogloss:0.54659
[97] train-mlogloss:0.31154 eval-mlogloss:0.54674
[98] train-mlogloss:0.31009 eval-mlogloss:0.54709
[99] train-mlogloss:0.30893 eval-mlogloss:0.54801
Accuracy: 0.8268194070080862

```

Classification Report:

	precision	recall	f1-score	support
0	0.67	0.09	0.15	23
1	0.17	0.02	0.03	220
2	0.84	0.98	0.91	1241
accuracy			0.83	1484
macro avg	0.56	0.36	0.36	1484
weighted avg	0.74	0.83	0.76	1484

Feature Importance

```
In [86]: importance = model.get_score(importance_type='weight')
importance_df = pd.DataFrame({
    'Feature': list(importance.keys()),
    'Importance': list(importance.values())
})

# Sort the DataFrame by importance
importance_df = importance_df.sort_values(by='Importance', ascending=False)

print('Feature Importances:')
print(importance_df)

# Plot feature importances
plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=importance_df)
plt.title('Feature Importances')
plt.show()
```

```
Feature Importances:
   Feature  Importance
2  Type_of_vehicle    2636.0
1   Day_of_week      2187.0
0    time_group      2182.0
4  Type_of_collision    1440.0
5   Casualty_class    1281.0
3  Road_surface_type     652.0
6   Sex_of_casualty     505.0
```

